

```

1 library(readr)
2 library(tidyverse)

— Attaching core tidyverse packages — tidyverse 2.0.0 —
✓ dplyr      1.1.4    ✓ purrr      1.0.2
✓ forcats    1.0.0    ✓ stringr    1.5.1
✓ ggplot2    3.4.4    ✓ tibble     3.2.1
✓ lubridate  1.9.3    ✓ tidyr      1.3.1
— Conflicts — tidyverse_conflicts() —
✖ dplyr::filter() masks stats::filter()
✖ dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

```

```

1 df <- read.csv('./Real estate.csv')
2
3 head(df)

```

A data.frame: 6 × 8

ansaction.date	X2.house.age	X3.distance.to.the.nearest.MRT.station	X4.number.of.conven
<dbl>	<dbl>		<dbl>
2012.917	32.0		84.87882
2012.917	19.5		306.59470
2013.583	13.3		561.98450
2013.500	13.3		561.98450
2012.833	5.0		390.56840
2012.667	7.1		2175.03000

Double-click (or enter) to edit

```

1 summary(df)

      No      X1.transaction.date X2.house.age
Min.   : 1.0   Min.   :2013         Min.    : 0.000
1st Qu.:104.2  1st Qu.:2013         1st Qu.: 9.025
Median :207.5  Median :2013         Median :16.100
Mean   :207.5  Mean   :2013         Mean   :17.713
3rd Qu.:310.8  3rd Qu.:2013         3rd Qu.:28.150
Max.   :414.0  Max.   :2014         Max.   :43.800
X3.distance.to.the.nearest.MRT.station X4.number.of.convenience.stores
Min.   : 23.38                        Min.    : 0.000
1st Qu.: 289.32                        1st Qu.: 1.000
Median : 492.23                        Median : 4.000
Mean   :1083.89                        Mean    : 4.094
3rd Qu.:1454.28                        3rd Qu.: 6.000
Max.   :6488.02                        Max.   :10.000
X5.latitude X6.longitude Y.house.price.of.unit.area
Min.   :24.93   Min.   :121.5   Min.    : 7.60
1st Qu.:24.96   1st Qu.:121.5   1st Qu.: 27.70
Median :24.97   Median :121.5   Median : 38.45
Mean   :24.97   Mean   :121.5   Mean    : 37.98
3rd Qu.:24.98   3rd Qu.:121.5   3rd Qu.: 46.60
Max.   :25.01   Max.   :121.6   Max.   :117.50

```

```

1 str(df)

'data.frame':   414 obs. of  8 variables:
 $ No           : int  1 2 3 4 5 6 7 8 9 10 ...
 $ X1.transaction.date : num  2013 2013 2014 2014 2013 ...
 $ X2.house.age   : num  32 19.5 13.3 13.3 5 7.1 34.5 20.3 31.7 17.9 ...
 $ X3.distance.to.the.nearest.MRT.station: num  84.9 306.6 562 562 390.6 ...
 $ X4.number.of.convenience.stores      : int  10 9 5 5 5 3 7 6 1 3 ...
 $ X5.latitude                          : num  25 25 25 25 25 ...
 $ X6.longitude                         : num  122 122 122 122 122 ...
 $ Y.house.price.of.unit.area           : num  37.9 42.2 47.3 54.8 43.1 32.1 40.3 46.7 18.8 22.1 ...

```

```

1 cor(df)

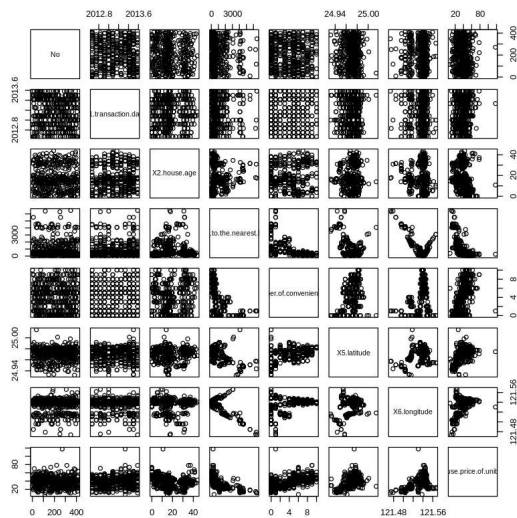
```

A matrix: 8 × 8 of type dbl

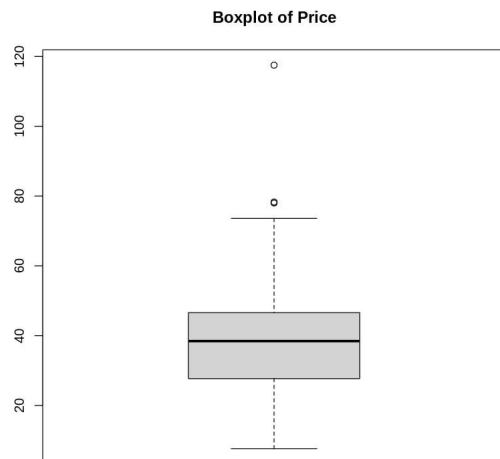
ansaction.date	X2.house.age	X3.distance.to.the.nearest.MRT.station	X4.number.of.conven
-0.048657949	-0.03280811		-0.01357349
1.000000000	0.01754877		0.06087995
0.017548767	1.00000000		0.02562205
0.060879953	0.02562205		1.00000000
0.009635445	0.04959251		-0.60251914
0.035057756	0.05441990		-0.59106657
-0.041081778	-0.04852005		-0.80631677
0.087490606	-0.21056705		-0.67361286

Double-click (or enter) to edit


1 pairs(df)



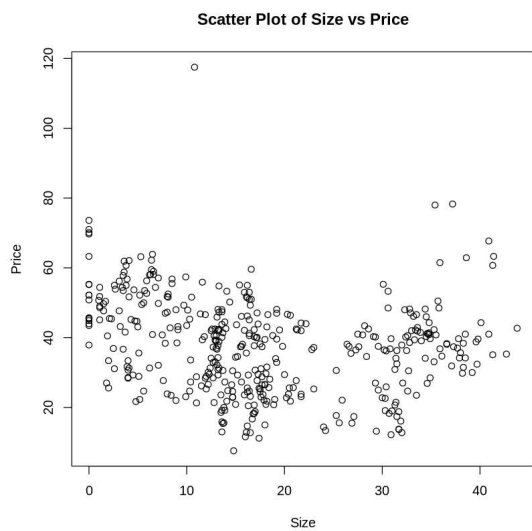
1 boxplot(df\$Y.house.price.of.unit.area, main="Boxplot of Price")



1 hist(df\$Y.house.price.of.unit.area, main="Histogram of Price", xlab="Price")



```
1 plot(df$X2.house.age, df$Y.house.price.of.unit.area, main="Scatter Plot of Size vs Price", xlab="Size", ylab="Price")
```



✓ .

.

.

.

.

.

.

.

.

.

```

1 library(ggplot2)
2 library(dplyr)
3 library(broom)

1 model <- lm(Y.house.price.of.unit.area ~ X1.transaction.date + X2.house.age +
2           X3.distance.to.the.nearest.MRT.station + X4.number.of.convenience.stores +
3           X5.latitude + X6.longitude, data = df)
4
5 summary(model)

```

Call:

```
lm(formula = Y.house.price.of.unit.area ~ X1.transaction.date +
    X2.house.age + X3.distance.to.the.nearest.MRT.station +
    X4.number.of.convenience.stores +
    X5.latitude + X6.longitude, data = df)
```

Residuals:

```

      Min       1Q   Median       3Q      Max
-35.664  -5.410  -0.966   4.217  75.193

```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.444e+04	6.776e+03	-2.131	0.03371
X1.transaction.date	5.146e+00	1.557e+00	3.305	0.00103
X2.house.age	-2.697e-01	3.853e-02	-7.000	1.06e-11
X3.distance.to.the.nearest.MRT.station	-4.488e-03	7.180e-04	-6.250	1.04e-09
X4.number.of.convenience.stores	1.133e+00	1.882e-01	6.023	3.84e-09
X5.latitude	2.255e+02	4.457e+01	5.059	6.38e-07
X6.longitude	-1.242e+01	4.858e+01	-0.256	0.79829

```

(Intercept)          *
X1.transaction.date   **
X2.house.age          ***
X3.distance.to.the.nearest.MRT.station ***
X4.number.of.convenience.stores ***
X5.latitude           ***
X6.longitude
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.858 on 407 degrees of freedom


Multiple R-squared: 0.5824, Adjusted R-squared: 0.5762

F-statistic: 94.59 on 6 and 407 DF, p-value: < 2.2e-16

```

1 # Plot the regression model
2 plot(df$Y.house.price.of.unit.area, fitted(model),
3      main="Regression Model", xlab="Observed Values", ylab="Fitted Values")
4
5

```



```

1 ggplot(df, aes(x = Y.house.price.of.unit.area, y = fitted(model))) +

```

```

1 ggplot(ut, aes(x = r.house.price.ot.unit.area, y = fitted(model))) +
2   geom_point() +
3   geom_smooth(method = "lm", se = FALSE, color = "red") +
4   ggtitle("Regression Model") +
5   xlab("Observed Values") +
6   ylab("Fitted Values") +
7   theme_minimal()
8

```

```

`geom_smooth()` using formula = 'y ~ x'
Regression Model

```

